

# Clustering Analysis in R

## An Introduction for DAT 280

Alex Marsella

February 22, 2024

# Today's gameplan

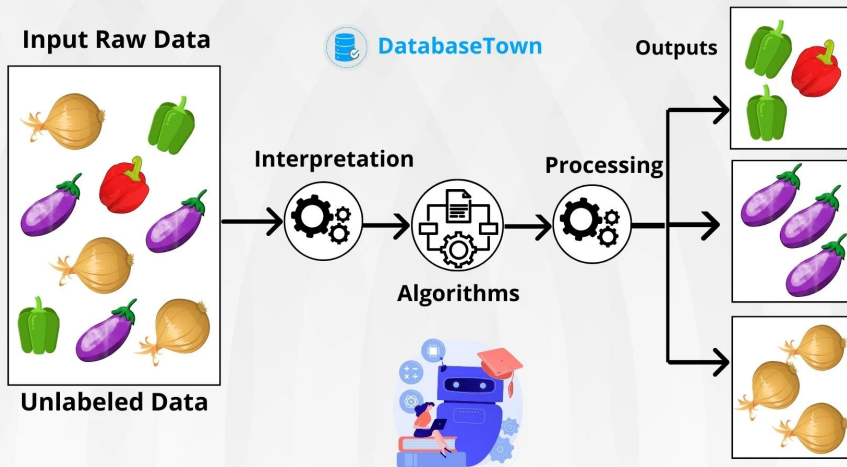
- I will introduce you to the idea of clustering, a form of machine learning.
- Today's "Lecture" lecture will be more lecture-y, while Monday's "Coding" lecture will be more hands-on.
- I still have an exercise at the end for you to try out.

# Introduction to Clustering

Clustering is a type of unsupervised learning method used to group similar entities together. It's a powerful tool in data analysis, allowing us to discover natural groupings in data based on their characteristics.

# UNSUPERVISED LEARNING

Unsupervised learning is a type of machine learning where the algorithm learns from unlabeled data without any predefined outputs or target variables.



# Real-World Applications of Clustering

Clustering has a wide range of applications across different fields:

- **Market segmentation:** Grouping customers based on purchasing behavior.
- **Biology:** Classifying plants or animals based on their features.
- **Image segmentation:** Dividing digital images into multiple segments (sets of pixels).
- **Anomaly detection:** Identifying unusual data points that do not fit into any cluster.

# Two types of clustering

- **Hierarchical Clustering:** Builds a hierarchy of clusters either through a bottom-up (agglomerative) or top-down (divisive) approach.
- **Partitioning Clustering:** Divides the data into non-overlapping subsets (clusters) such that each data point is in exactly one subset. The most common method is k-means clustering.

# K-means Clustering

K-means clustering groups  $n$  observations into  $k$  clusters in which each observation belongs to the cluster with the nearest mean.

## Algorithm Steps

- 1 **Initialization:** Start with  $k$  centroids by *randomly* selecting  $k$  points from the dataset.
- 2 **Assignment:** Assign each point to the nearest centroid.
- 3 **Update:** Recalculate the centroids as the center of the points assigned to each cluster.
- 4 **Repeat:** Repeat steps 2 and 3 until the centroids no longer change significantly.

Before we move any further, let's load in tidyverse, cluster, and iris

```
library(tidyverse)
library(cluster)
data(iris)
```



# K-means clustering in the iris dataset

- We will use the `kmeans()` command on columns 1:4 of `iris` and tell it to make 3 clusters.
- This will assign all irises to one of three “clusters” based on the characteristics in those four columns.
  - Sepal width and length, petal width and length
- Note that we do NOT cluster on the categorical aspect of species.
  - Species is nature’s clustering, so to speak, we want to see if we can predict that or come up with our own clustering.


```
set.seed(123)  # for reproducibility
kmeans_result <- kmeans(iris[, 1:4], centers = 3)
```

## Plotting the result.

- We color it by the clusters stored in `kmeans_result` even though we are using `iris`
  - note that I wrap it in `factor()` since the clusters are stored as numeric but I want them to be read as categories.

```
ggplot(iris, aes(Sepal.Length, Sepal.Width,  
  color = factor(kmeans_result$cluster))) +  
  geom_point() + labs(title = "K-means Clustering of the Iris",  
  color = "Cluster")
```


### K-means Clustering of the Iris Dataset



Plotting the result, again.

```
ggplot(iris, aes(Petal.Length, Petal.Width,  
  color = factor(kmeans_result$cluster))) +  
  geom_point() + labs(title = "K-means Clustering of the Iris  
  color = "Cluster")
```


## K-means Clustering of the Iris Dataset



## Comparing it to the three species of Iris

- Clustering can provide prediction about a class to which something belongs.
- Notice how the clusters are fairly congruent with the species of Iris.

```
ggplot(iris, aes(Sepal.Length, Sepal.Width,  
  color = factor(kmeans_result$cluster))) +  
  geom_point(aes(shape = Species)) +  
  labs(title = "K-means Clustering of Iris Dataset",  
    color = "Cluster")
```



# Hierarchical Clustering

Hierarchical clustering creates a tree of clusters. It doesn't require us to pre-specify the number of clusters. Instead, it produces a dendrogram, allowing us to choose the number of clusters based on the tree.


## Agglomerative Hierarchical Clustering

- 1 **Start:** Treats each point as a separate cluster.
- 2 **Find:** Identifies the closest two clusters and merges them.
- 3 **Repeat:** Continue the process until all points are clustered into a single group.
- 4 **Dendrogram:** A tree-like diagram that records the sequences of merges or splits.

# Hierarchical Clustering

```
dist_mat <- dist(iris[, 1:4]) # Calculate distance matrix
hc <- hclust(dist_mat) # Perform hierarchical clustering
plot(hc, main = "Hierarchical Clustering Dendrogram",
      xlab = "", sub = "") # Plot the dendrogram
```

Hierarchical Clustering Dendrogram



# Choosing the right number of clusters


## The Elbow Method

The Elbow Method is a popular technique to determine the optimal number of clusters  $k$  in K-means clustering. It involves the following steps:

- 1 **Compute Clustering:** Perform K-means clustering on the dataset for a range of values of  $k$  (e.g., 1 to 10).
- 2 **Calculate Within-Cluster Sum of Square (WCSS):** For each  $k$ , calculate the total within-cluster sum of square (WCSS).
- 3 **Plot the Curve:** Plot  $k$  against the WCSS. The plot typically shows a rapid decline in WCSS as  $k$  increases, which eventually slows, creating an “elbow”.
- 4 **Determine the Elbow Point:** The point where the rate of decrease sharply changes (the elbow) represents the optimal  $k$ .

# The “Within-Cluster Sum of Squares”


- Sum of the squared distances between an observation in a cluster and its centroid.



$$WCSS = \sum_{P_i \text{ in Cluster 1}} \text{distance}(P_i, C_1)^2 + \sum_{P_i \text{ in Cluster 2}} \text{distance}(P_i, C_2)^2 + \sum_{P_i \text{ in Cluster 3}} \text{distance}(P_i, C_3)^2$$



$$\sqrt{(X_1 - X_2)^2 + (Y_1 - Y_2)^2}$$



$$\sqrt{(A_1 - A_2)^2 + (B_1 - B_2)^2 + \dots + (Z_1 - Z_2)^2}$$

# Why Use the Elbow Method?

- **Simplicity:** Easy to understand and implement.
- **Heuristic:** Helps in making an informed decision based on the WCSS plot.

## Limitations

- **Subjectivity:** The exact “elbow” point can sometimes be subjective or not very clear.
- **Not Always Accurate:** May not always provide the best number of clusters for complex datasets.

## Coding the setup

- Remove the species column (fifth column) to only include features.
- Create a vector to store WCSS values for 1-10 clusters.
- Create a “for loop” in R to calculate WCSS 10 times for 1-10 clusters.


```
iris_data <- iris[, -5]

wcss <- numeric(10)

for (k in 1:10) {
  # k will take on the value 1,
  # then 2, and so on, repeating
  # to 10
  kmeans_result <- kmeans(iris_data,
    centers = k, nstart = 20)
  wcss[k] <- kmeans_result$tot.withinss
}
```

## Creating the plot: Is two clusters enough for Irises?

```
ggplot(, aes(x = 1:10, y = wcss)) + geom_line() +  
  geom_point() + xlab("Number of Clusters k") +  
  ylab("Within-Cluster Sum of Squares (WCSS)") +  
  theme_minimal()
```



# Assessing Cluster Quality

Understanding the quality of the clusters formed is crucial for validating the results of your clustering analysis.

## Silhouette Analysis


- Silhouette analysis measures how similar an object is to its own cluster compared to other clusters.
- The silhouette score ranges from -1 to 1, where a high value indicates that the object is well matched to its own cluster and poorly matched to neighboring clusters.
  - $s > 0.7$  “strong”
  - $0.7 > s > 0.5$  “reasonable”
  - $0.5 > s > 0.25$  “weak”
  - Close to 0 implies overlapping clusters
  - Negative values imply incorrect assignment to clusters.

# Silhouette Analysis

```
sil <- silhouette(kmeans_result$cluster,  
  dist(iris[, 1:4]))  
plot(sil, main = "Silhouette Analysis") # base R plotting is
```

## Silhouette Analysis

n = 150



# Summary & Final Thoughts

- Clustering is a versatile tool in data analysis, with numerous applications.
- The choice of algorithm depends on the dataset and the specific requirements of the analysis.
- Assessing cluster quality is essential for ensuring meaningful results.
- Continuous exploration and learning are key, as new clustering methods and applications are regularly developed.

## Exercise

Apply clustering analysis to segment customers based on their purchasing behavior. This is a common task in marketing and business strategy, helping companies tailor their approaches to different customer groups.

### Dataset Description

The dataset, `CustomerData.csv`, represents purchasing data collected from a retail store. It includes the following variables:

- `CustomerID`: Unique identifier for each customer.
- `AnnualIncome`: The annual income of the customer (in thousands).
- `SpendingScore`: A score assigned by the mall based on customer behavior and spending nature (1-100).



# Objective

Your task is to segment the customers into distinct groups based on their `AnnualIncome` and `SpendingScore`. You will:

- 1 Perform exploratory data analysis (EDA) to understand the dataset.
- 2 Use k-means clustering to identify customer segments.
- 3 Evaluate the clustering and interpret the customer segments.

## Step 1: Load the Dataset

Start by loading the dataset into R and taking a peek. I have set this up for you to expedite it.

```
customer_data <- read.csv("CustomerData.csv")  
head(customer_data)
```

##	CustomerID	AnnualIncome	SpendingScore
## 1	1	54	25
## 2	2	121	96
## 3	3	70	61
## 4	4	134	52
## 5	5	142	41
## 6	6	21	88

## Step 2: Exploratory Data Analysis (EDA)

I have already performed a basic exploratory analysis for you.


```
# Summary statistics  
summary(customer_data[, c("AnnualIncome",  
    "SpendingScore")])
```

##	AnnualIncome	SpendingScore
##	Min. : 15.00	Min. : 2.00
##	1st Qu.: 51.75	1st Qu.: 24.75
##	Median : 80.00	Median : 47.50
##	Mean : 83.34	Mean : 49.44
##	3rd Qu.: 114.25	3rd Qu.: 74.50
##	Max. : 149.00	Max. : 100.00

## Step 2: Exploratory Data Analysis (EDA)

- Notice how there is no discernible trend at all.

```
ggplot(customer_data, aes(x = AnnualIncome,  
  y = SpendingScore)) + geom_point() +  
  labs(title = "Annual Income vs. Spending Score",  
    x = "Annual Income (in thousands)",  
    y = "Spending Score") + geom_smooth(method = "lm")
```



## Step 3: Find the elbow

Use the elbow method to determine a good number of clusters ( $k$ ).

## Step 4: K-means Clustering

Use k-means clustering to segment the customers. Use four clusters.

## Step 5: Visualize it

- Plot the results with a unique trendline for each of the four customer segments.

## Step 5: Describe and Interpret

How does the relationship between Annual Income and Spending Score differ by customer type?